

## Part II: R

I. Create a linear regression analysis with R to predict the size of the population for the state you live in for 2020 based on the “Current Estimates Data” dataset (see weblink below). Provide a screenshot of your results [2].

```
```${r echo=FALSE, message=FALSE, warning=FALSE}
#Confidence Interval for the Utah population estimate for the year 2020
y_pred <- lm(regressor, data = test_set)
new_data <- data.frame(Year=2020)
utah_predict <- predict(y_pred, new_data, interval="predict", level=.95)
utah_predict
```
```

|   | fit     | lwr     | upr     |
|---|---------|---------|---------|
| 1 | 3254141 | 3002358 | 3505925 |

Using a 95% confidence interval, the results of the linear regression model indicate the estimated population for the state of Utah will be between 3,002,358 and 3,505,925 people.

```
```${r echo=FALSE, message=FALSE, warning=FALSE}
#Utah state population estimate for the year 2020
utah_predict <- predict(y_pred, new_data)
utah_predict
```
```

|         |
|---------|
| 1       |
| 3254141 |

The Linear Regression model result for the 2020 Utah state population is 3,254,141 people.

**J. Explain how you prepared the data from part I and how the dataset was imported into R, including screenshots of your results.**

Step 1: Downloaded the "Current Estimates" dataset as a CSV file. Using MS Excel, I deleted all unnecessary columns, keeping only the Utah census data from 2010 – 2017. I then saved the cleaned dataset as a new CSV file named ut\_census.

Step 2: Using the Readr library, I imported the cleaned Utah Census Bureau data into R.

```
``{r echo=FALSE, message=FALSE, warning=FALSE}
# Load the Data
UT_data<- read_csv("~/Desktop/MSDA Portfolio/2. DS Tools &
Techniques/P2. R/ut_census.csv")

View(UT_data)
```

|   | Year | Pop_Est |
|---|------|---------|
| 1 | 2010 | 2775260 |
| 2 | 2011 | 2815430 |
| 3 | 2012 | 2854222 |
| 4 | 2013 | 2899961 |
| 5 | 2014 | 2938671 |
| 6 | 2015 | 2984917 |
| 7 | 2016 | 3044321 |
| 8 | 2017 | 3101833 |

K. Create an R script that will tabulate a statistical description of the estimated 2020 data using R's `summary()` function and provide a screenshot of your results. [2]

```
```{r message=FALSE, warning=FALSE}
# Split the dataset into the Training set and Test set
set.seed(123)
split = sample.split(UT_data$Pop_Est, SplitRatio = 2/3)
training_set = subset(UT_data, split == TRUE)
test_set = subset(UT_data, split == FALSE)

# Fit the Simple Linear Regression to the Training set
regressor = lm(formula = Pop_Est ~ Year,
               data = training_set)
```
```

```
```{r echo=FALSE, message=FALSE, warning=FALSE}
summary(regressor)
```
```

**Call:**  
`lm(formula = Pop_Est ~ Year, data = training_set)`

**Residuals:**

| 1      | 2     | 3       | 4       | 5      |
|--------|-------|---------|---------|--------|
| 4420.4 | 308.1 | -5182.2 | -7334.0 | 7787.7 |

**Coefficients:**

|             | Estimate  | Std. Error | t value | Pr(> t ) |     |
|-------------|-----------|------------|---------|----------|-----|
| (Intercept) | -86236535 | 2847644    | -30.28  | 7.91e-05 | *** |
| Year        | 44282     | 1415       | 31.30   | 7.17e-05 | *** |

---

**Signif. codes:** 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

**Residual standard error:** 7324 on 3 degrees of freedom  
**Multiple R-squared:** 0.9969, **Adjusted R-squared:** 0.9959  
**F-statistic:** 979.7 on 1 and 3 DF, **p-value:** 7.165e-05

L. Predict the population size of your state in five years using a linear regression from part I and provide a screenshot of your results [2].

```
```{r echo=FALSE, message=FALSE, warning=FALSE}
#Confidence Interval for the Utah population estimate for the year 2023
y_pred <- lm(regressor, data = test_set)
new_data <-data.frame(Year=2023)
utah_predict<- predict(y_pred,new_data,interval="predict",level=.95)
utah_predict
```
```

|   | fit     | lwr     | upr     |
|---|---------|---------|---------|
| 1 | 3408259 | 3049965 | 3766552 |

Using a 95% confidence interval, the results of the linear regression model indicate the estimated population for the state of Utah will be between 3,049,965 and 3,766,552 people.

```
```{r echo=FALSE, message=FALSE, warning=FALSE}
#Utah state population estimate for the year 2023
utah_predict<- predict(y_pred,new_data)
utah_predict
```
```

|         |
|---------|
| 1       |
| 3408259 |

Predicted Utah state population for the year 2023 is 3,408,259 people.

## Resources

1. Dataset - Retrieved August 30, 2018, from <https://www2.census.gov/programs-surveys/popest/datasets/2010-2017/state/asrh/>
2. Eremenko, K., & De Ponteves, H. (n.d.). Machine Learning A-Z™: Hands-On Python & R In Data Science. Retrieved September 1, 2018, from <https://www.udemy.com/>